

AMR Prediction: Lessons from Distribution Shift

MohammadErfan Jabbari^{1,2} & Ana-Maria Mirza¹

¹Universidad Carlos III de Madrid · ²IMDEA Networks Institute

1. The Challenges

Three properties of this dataset shaped every modeling decision:

1. Species Distribution Shift

The training set is dominated by *P. aeruginosa* (43%), but this species represents only 3% of the test set. Meanwhile, *K. pneumoniae* jumps from 28% to 51%. Any model that performs well on training validation will be optimizing for the wrong population.




Figure 1: *P. aeruginosa* dominates training (43%) but nearly vanishes in test (3%). Models optimizing for training distribution will fail.

2. Extreme Feature Sparsity

93% of the MALDI-TOF feature matrix is zeros. The signal lives in sparse intensity peaks, not dense patterns. This favors tree-based models (LightGBM) over neural networks, and makes aggressive dimensionality reduction dangerous: compress too hard and you discard the few features that matter.

3. Non-Random Missing Labels

Label missingness is not uniform: 43% of Amoxicillin/Clavulanic acid labels are missing, compared to under 3% for most other antibiotics. This reflects clinical testing practices, not random noise. Ignoring the pattern or imputing naively corrupts the learning signal.

These three facts explain most of what worked and what failed.

2. What We Explored

We systematically tested approaches across five categories before converging on the final pipeline:



Figure 2: Failures were informative: stacking overfit, species-specific models lacked data, aggressive DR lost signal.

3. Results

Submission	Public LB	Private LB	OOF CV
Mega-blend (rank avg)	0.8386	0.8131	0.8174
Self-training blend	0.8366	0.8094	0.8052
LightGBM + PLS	0.8323	0.8012	0.7986
Stacking (overfit)	0.8269	0.7840	0.9545
Baseline	0.8022	0.7856	0.7823

Table 1: Stacking achieved OOF 0.9545 but collapsed on private leaderboard.

Lesson: Under distribution shift, simple and robust beats clever and brittle.

4. What Worked

Three techniques lifted our score from 0.80 baseline to 0.81 private:

1. Rank Averaging over Stacking

Stacking learns to trust your validation signal. Under distribution shift, that trust is misplaced. Rank averaging converts predictions to ranks before blending, normalizing scale differences between models and reducing sensitivity to outliers. It stays skeptical.

$$\hat{y}_{\text{blend}} = \frac{1}{M} \sum_{m=1}^M \frac{\text{rank}(\hat{y}_m)}{N} \quad (\text{M=3 models, N=samples})$$

2. Species-Aware Weighting

Knowing that *K. pneumoniae* dominates the test set (51%) and *P. aeruginosa* nearly vanishes (3%), we up-weighted *K. pneumoniae* predictions by $3\times$ and down-weighted *P. aeruginosa* to $0.1\times$. Simple rebalancing that matches the deployment distribution.

3. Intrinsic Resistance Rules

Certain species are intrinsically resistant to certain antibiotics. This biological prior gives us 18.8% of predictions for free with 100% confidence. We override model outputs with hard-coded resistance where clinical rules apply.




Figure 3: Final pipeline. Training: 5-fold stratified CV, LightGBM (lr=0.03, 300 trees, 31 leaves), PLS 100 components.

Why PLS? Partial Least Squares projects features using target labels, preserving signal that correlates with resistance. Unsupervised methods (PCA, NMF) optimize variance, discarding predictive but low-variance peaks.

5. What Failed & Why

Honest lessons from approaches that did not improve our score:

Stacking (OOF 0.9545 → LB 0.827)

The meta-learner memorized validation patterns that did not generalize. High OOF score was a trap, not a signal.

Species-Specific Models (32 separate models)

Splitting by species left too few samples per model. *P. aeruginosa* had 4,100 training rows but represents only 3% of test. Underfitting dominated.

Aggressive Dimensionality Reduction

KPCA and NMF compressed too hard, discarding sparse intensity peaks that carried signal. PLS succeeded because it uses the target during projection.

K. pneumoniae-Only Optimization

We tried tuning exclusively for *K. pneumoniae* (51% of test). This hurt *E. coli* and *P. mirabilis* enough to lower overall score. The metric is macro-averaged.

*“Stacking learns to trust your validation signal.
Rank averaging stays skeptical.
Under distribution shift, skepticism wins.”*

6. References

- [1] Weis et al. *Nature Medicine* 28, 164–174 (2022). DRIAMS dataset.
- [2] Ke et al. *NeurIPS* 30, 3146–3154 (2017). LightGBM.
- [3] Chen & Guestrin. *KDD*, 785–794 (2016). XGBoost.
- [4] Prokhorenkova et al. *NeurIPS* 31, 6638–6648 (2018). CatBoost.